

EDUCATION

Ph.D. Student to M.S., AI for Geosciences University of Texas at Austin | 2024 – Expected Spring 2027
B.S. in Geophysics (Honors) University of Science and Technology of China | 2020 – 2024
• Nominated for the **Highest honor** for USTC undergrads (one of the two ESS School nominees) Mar. 2024

AI EVALUATION EXPERIENCE

Tencent Hunyuan AI Evaluation Qingyun Offer Declined (University Administrative Constraints) | Aug. 2026
• Gathered insights through AI events and direct conversations with researchers and industry participants, informing which benchmarks/data source opportunities mattered to the [Hunyuan](#) multi-modal foundation model team.
• Pre-intern Evaluation Project (one month): Discussed with data vendors and analyzed 40 leading benchmarks and 300+ unique profiles to identify potential data-sourcing and collaboration partners. Checked the quality of vendor data samples, delivering reproducible findings and built a reusable pipeline for STEM sample evaluations.

Creator, Benchmark Radar Jul. 2026 – Present
• Built an open-source [live radar](#) to make fragmented benchmark, evaluation, and dataset work discoverable early; collected and classified 6k+ signals, getting 100+ stars and 20+ forks within the first month. [\[Code\]](#)
• Attracted 100k+ media views and attention from [Google](#), [NVIDIA](#), [Kimi](#), [Qwen](#), and [Jiniqu Capital](#). [\[Posts\]](#)

Research Intern, Production Agent Benchmarking [19Pine.AI](#) (Singapore) | Remote | Jul. 2025 – Aug. 2025
• Built a prototype of PineAgent's production benchmarking system for a personal agent handling high-stakes phone tasks such as customer claims, evaluating RAG accuracy, privacy safety, and task completion based on 1,000+ call sessions. *Pine later ranked No. 1 on the τ -Voice leaderboard, outperforming GPT Realtime and Gemini Live.*

Lead Developer, ESM-bench for Earth System Models UT-Austin | Dec. 2025 – Present
• Created 107-task ESM-bench to test physical validity and Fortran syntax correctness, which ordinary coding benchmarks miss; <20% exact match exposed a core weakness in frontier coding agents. [\[Preprint\]](#)

Core Contributor, ASI-Bench Apex Intelligence | Apr. 2026 – Jul. 2026
• Authored 3 Earth Science tasks and strengthened benchmark infra on task and scorer validity. [\[arXiv\]](#)

ResearchClawBench Contributor: Noah-MP science task. [\[arXiv\]](#) [\[Hugging Face #2 Paper of the day\]](#)
Open-Source Contributor, AI Safety Audit Infrastructure [meridianlabs-ai/inspect_petri](#) | Jun. 2026 – Present
• Landed 2 tested upstream fixes in Petri, an alignment-auditing agent whose transcripts serve as safety evidence: judge-visible auditor prefill attribution and shipped benchmark-seed canary coverage. [\[PR\]](#)

Research Collaborator, MatraIx Global Population-Scale Simulated-User Evaluation | Jul. 2026 – Present
• Built a resumable Qwen3.6 vLLM pipeline orchestrating 1.6K GPU-hours to extract 38k unique personas with zero cross-system overlap. Coauthor, 11th of 93 authors for *MatraIx: Simulating the World with 8.3 Billion Persona Agents* featured by [Forbes Türkiye](#) and [WIRED Czech](#); [Hugging Face Daily Papers](#).

PUBLICATIONS

Selected from 10 papers, preprints, and chapters; see [Google Scholar](#).

- Lead author, *ESM-Bench: Evaluating Whether AI Agents Understand Earth System Model Physics and Code*
- Lead author, *Noah-Agent: A Multi-Expert Agent Autoresearch Framework for Large-Scale Climate Models*
- First author, *JGR: Space Physics* (2024).
- Co-author (11th of 93), *MatraIx: Simulating the World with 8.3 Billion Persona Agents* [\[HF Daily Papers\]](#)
- Co-author, *ResearchClawBench: End-to-End Autonomous Scientific Research*. [\[Hugging Face #2 Paper of the day\]](#)
- Co-author, *ASI-Bench: At the Dawn of Artificial Superintelligence*. [\[Website\]](#) [\[Code\]](#)

TECHNICAL SKILLS

- **Evaluation:** benchmark and dataset construction, rubric design, LLM-as-judge, task validators, hidden tests, failure taxonomies, red-teaming, prompt injection and adversarial testing, experiment analysis, evaluation reporting.
- **Engineering:** Python, PyTorch, NumPy, pandas, xarray, SQL, TypeScript/JavaScript, FastAPI, Node.js, Docker, GitHub Actions, AWS/GCP, Linux/HPC.

STARTUP EXPERIENCE

Technical Co-Founder, Tacite.AI Austin, TX | Jan. 2026 – Present
• Designed and built two proactive AI assistants, Tacite APP for individuals and Tacite COS for teams, on a local-first React/Electron/SQLite architecture that keeps user data on-device and supports background scheduling.
• Used test-driven development with hundreds of passing tests and shipped cross-platform macOS and Windows builds.
• Shaped product direction around privacy, responsiveness, local-first workflows, and user control. [\[Website\]](#)